


Storage Management stuff

Sukhjit Singh

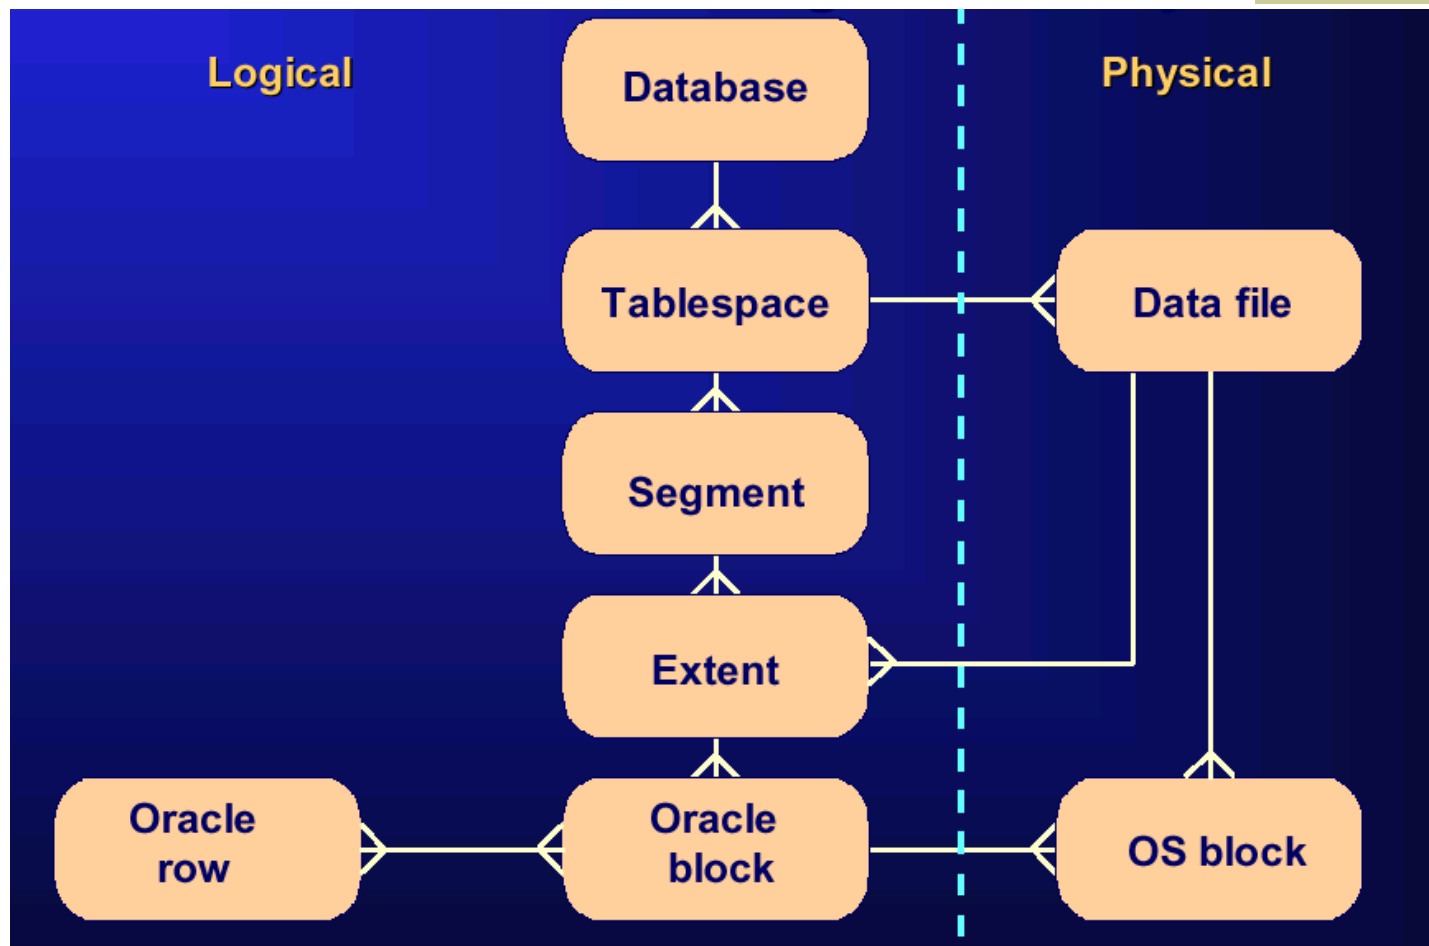
Agenda

- ◆ Data Storage Clause – including Storage Mgmt
- ◆ Sequences, Synonyms
- ◆ Indexes

Storage Management



Data Storage and Management



Segment Organization

- **Application segments**
 - **With data: Table, cluster, table partition, nested table, LOB, IOT, IOT overflow**
 - **Used to access data: Index, index partition, IOT, LOB index**
- **System segments: Rollback, temporary, bootstrap**

Storage Management Parameters

The following parameters influence the segment storage allocation:

- **INITIAL:** Size of the first extent in bytes
- **NEXT:** Starting size of subsequent extents
- **PCTINCREASE:** Percent increase of third and subsequent extents' size
- **MINEXTENTS:** Extents initially allocated
- **MAXEXTENTS:** Maximum extents allowed for a segment


Storage Mgmt. – Example

```
CREATE TABLE emps (  
    empno NUMBER(4), name VARCHAR2(30),  
    deptno NUMBER(3), hire_date DATE )  
STORAGE  
    ( INITIAL          2M  
      NEXT             500K  
      PCTINCREASE      0  
      MINEXTENTS       1  
      MAXEXTENTS       50 )  
TABLESPACE case_large_data;
```

Storage Clause


- ◆ Can be specified at Tablespace level and Table Level
- ◆ Default Storage Clause in “Create TableSpace” command, sets the default properties for Tables
- ◆ Blocks, Extents are allocated when a Segment is allocated.

DataBase Blocks Organization



- ◆ Consists of one or more data blocks organized as follows.

Storage Clause Usage



Storage Clause – Dictionaries

◆ • Data Dictionary Views

- – DBA_TABLESPACES
- – DBA_DATA_FILES
- – DBA_SEGMENTS
- – DBA_EXTENTS
- – DBA_FREE_SPACE



SEQUENCES

SEQUENCES

- Automatically generates unique numbers
- Is a sharable object
- Is typically used to create a primary key value
- Replaces application code
- Speeds up the efficiency of accessing sequence values when cached in memory

Create Sequence

- ◆ Define a sequence to generate sequential numbers automatically.

```
CREATE SEQUENCE sequence
    [INCREMENT BY n]
    [START WITH n]
    [{MAXVALUE n | NOMAXVALUE}]
    [{MINVALUE n | NOMINVALUE}]
    [{CYCLE | NOCYCLE}]
    [{CACHE n | NOCACHE}];
```

Create Sequence - Example

- Create a sequence named DEPT_DEPTNO to be used for the primary key of the DEPT table.
- Do not use the CYCLE option.

```
SQL> CREATE SEQUENCE dept_deptno  
2          INCREMENT BY 1  
3          START WITH 91  
4          MAXVALUE 100  
5          NOCACHE  
6          NOCYCLE;
```

Sequence created.

Sequences – Dictionary View

- Verify your sequence values in the USER_SEQUENCES data dictionary table.

```
SQL> SELECT  sequence_name, min_value, max_value,  
2           increment_by, last_number  
3 FROM      user_sequences;
```

- The LAST_NUMBER column displays the next available sequence number.

NEXTVAL OR CURRVAL

- ◆ NEXTVAL returns the next available sequence value.

It returns a unique value every time it is referenced, even for different users.

- ◆ CURRVAL obtains the current sequence value.

NEXTVAL must be issued for that sequence before CURRVAL contains a value.

Using a Sequence

- Insert a new department named “MARKETING” in San Diego.

```
SQL> INSERT INTO      dept(deptno, dname, loc)
      2  VALUES      (dept_deptno.NEXTVAL,
      3                'MARKETING', 'SAN DIEGO');
1 row created.
```

- View the current value for the DEPT_DEPTNO sequence.

```
SQL> SELECT      dept_deptno.CURRVAL
      2  FROM      dual;
```

Modify Sequence

```
SQL> ALTER SEQUENCE dept_deptno  
2      INCREMENT BY 1  
3      MAXVALUE 999999  
4      NOCACHE  
5      NOCYCLE;
```

Sequence altered.

Using a Sequence

- ◆ Caching sequence values in memory allows faster access to those values.
- ◆ Gaps in sequence values can occur when:
 - A rollback occurs
 - The system crashes
 - A sequence is used in another table
- ◆ View the next available sequence, if it was created with NOCACHE, by querying the USER_SEQUENCES table.
- ◆ Change the increment value, maximum value, minimum value, cycle option, or cache option.

Modify Sequence

- ◆ Must be the owner or have the ALTER privilege for the sequence.
- ◆ Only future sequence numbers are affected.
- ◆ The sequence must be dropped and re-created to restart the sequence at a different number.
- ◆ Some validation is performed.

Dropping Sequence

- ◆ Remove a sequence from the data dictionary by using the DROP SEQUENCE statement.
- ◆ Once removed, the sequence can no longer be referenced.

```
SQL> DROP SEQUENCE dept_deptno;  
Sequence dropped.
```



Synonyms

Synonyms

- ◆ Simplify access to objects by creating a synonym (another name for an object).
 - Refer to a table owned by another user.
 - Shorten lengthy object names.

```
CREATE [PUBLIC] SYNONYM synonym  
FOR      object;
```


Synonyms

- ◆ Create a shortened name for the DEPT_SUM_VU view.

```
SQL> CREATE SYNONYM d_sum  
2 FOR dept_sum_vu;  
Synonym Created.
```

- ◆ Drop a Synonyms

```
SQL> DROP SYNONYM d_sum;  
Synonym dropped.
```



Indexes

Index

- ◆ Is a schema object
- ◆ Is used by the Oracle Server to speed up the retrieval of rows by using a pointer
- ◆ Can reduce disk I/O by using rapid path access method to locate the data quickly
- ◆ Is independent of the table it indexes
- ◆ Is used and maintained automatically by the Oracle Server

Index

- ◆ Created automatically, when you define a **PRIMARY KEY** or **UNIQUE** constraint in a table definition.
- ◆ Created manually: Users can create nonunique indexes on columns to speed up access time to the rows.

Index

- ◆ An Index can be created one or more columns – using a composite key.

```
CREATE INDEX index  
ON table (column[, column]...);
```

- ◆ Improve the speed of query access on the ENAME column in the EMP table.

```
SQL> CREATE INDEX      emp_ename_idx  
      2  ON              emp(ename);  
Index created.
```

When to create Indexes?

- ◆ The column is used frequently in the WHERE clause or in a join condition.
- ◆ The column contains a wide range of values.
- ◆ The column contains a large number of null values.
- ◆ Two or more columns are frequently used together in a WHERE clause or a join condition.
- ◆ The table is large and most queries are expected to retrieve less than 2–4% of the rows.

When not to create Indexes?

- ◆ The table is small.
- ◆ The columns are not often used as a condition in the query.
- ◆ Most queries are expected to retrieve more than 2–4% of the rows.
- ◆ The table is updated frequently.

Index Dictionaries

- ◆ The USER_INDEXES data dictionary view contains the name of the index and its uniqueness.
- ◆ The USER_IND_COLUMNS view contains the index name, the table name, and the column name.

```
SQL> SELECT  ic.index_name, ic.column_name,  
2           ic.column_position col_pos, ix.uniqueness  
3 FROM      user_indexes ix, user_ind_columns ic  
4 WHERE     ic.index_name = ix.index_name  
5 AND       ic.table_name = 'EMP';
```


Function Based Indexes

- ◆ A function-based index is an index based on expressions.
- ◆ The index expression is built from table columns, constants, SQL functions, and user-defined functions.

```
SQL> CREATE TABLE test (col1 NUMBER);
```

```
SQL> CREATE INDEX test_index on test(col1,col1+10);
```

```
SQL> SELECT col1+10 FROM test;
```

Remove Index

- ◆ Remove an index from the data dictionary.

```
SQL> DROP INDEX index;
```

- ◆ Remove the EMP_ENAME_IDX index from the data dictionary.

```
SQL> DROP INDEX emp_ename_idx;  
Index dropped.
```

- ◆ To drop an index, you must be the owner of the index or have the DROP ANY INDEX privilege.

Unique vs Non-Unique Indexes

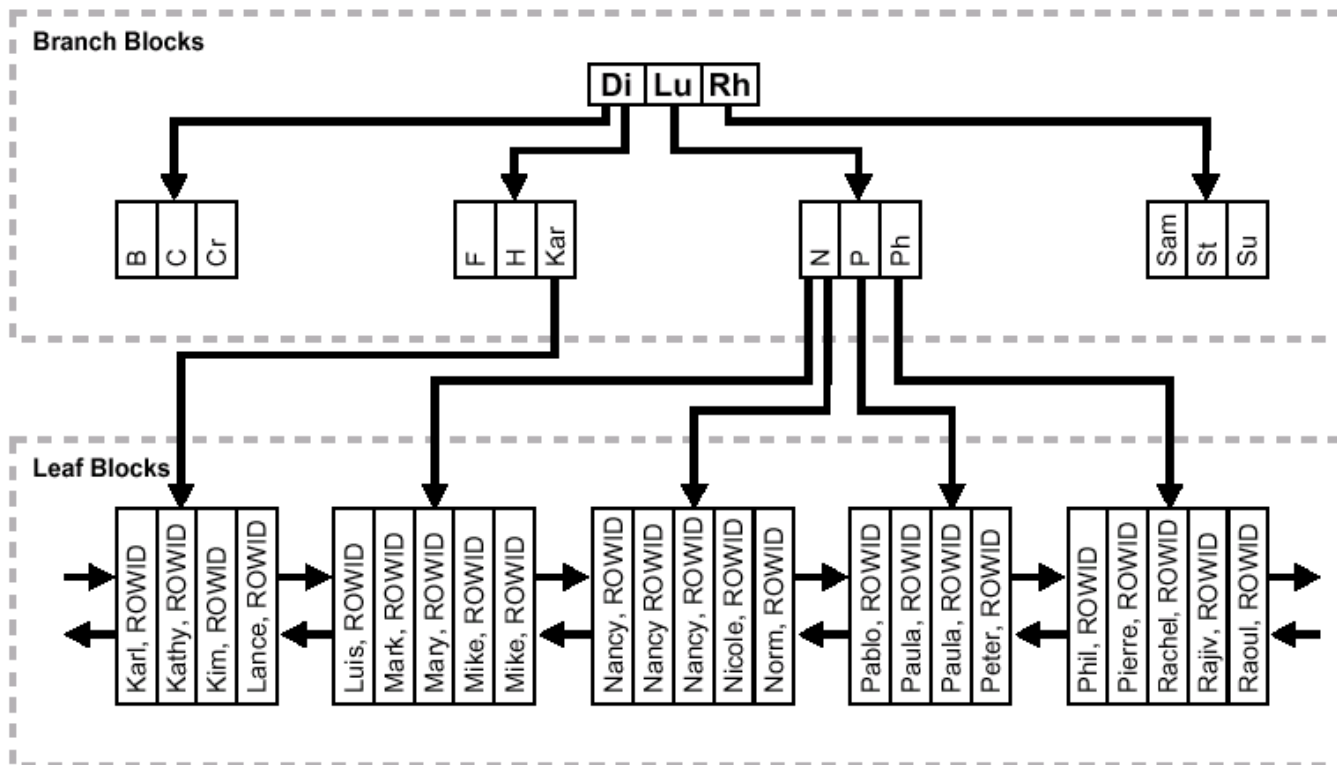
- ◆ Unique Indexes are created with Primary Keys. Assure no duplicates records are stored in table.
- ◆ Non-Unique Indexes impose no restriction on data columns.
- ◆ Unique Indexes can also be created explicitly on columns that unique constraint on them.
- ◆ NULL Values with Unique Indexes are considered to be distinct.

Index Storage

- ◆ B-tree indexes
- ◆ Bitmap indexes
- ◆ Bitmap Join indexes
- ◆ Indexes useful in clustering (Covered in future classes.)
 - B-tree cluster indexes
 - Hash cluster indexes
 - Reverse key indexes

Index Storage

◆ Balanced B-Tree



Index Storage

◆ Bitmap Index –

- Columns with low cardinality are good as Bitmap Indexes.
- Provide Good performance with And or OR Operations.
- Columns with high cardinality are not good candidates for Bitmap indexes.
- Bitmap Indexes include columns that have NULLS.

Index Storage

◆ Bitmap Index Example

CUSTOMER #	MARITAL_STATUS	REGION	GENDER	INCOME_LEVEL
101	single	east	male	bracket_1
102	married	central	female	bracket_4
103	married	west	female	bracket_2
104	divorced	west	male	bracket_4
105	s	REGION='east' REGION='central' REGION='west'		
106	r	1	0	0
		0	1	0
		0	0	1
		0	0	1
		0	1	0
		0	1	0

status = 'married'	region = 'central'	region = 'west'
0	0	0
1	1	0
1	0	1
0	0	1
0	1	0
1	1	0

status = 'married'	region = 'central'	region = 'west'
0	0	0
1	1	0
1	0	1
0	0	1
0	1	0
1	1	0

status = 'married'	region = 'central'	region = 'west'
0	0	0
1	1	0
1	0	1
0	0	1
0	1	0
1	1	0

Index Storage

◆ Bitmap Join Index

- A bitmap index created using two or more tables with a join.
- Used Traditionally in Data-Warehousing.
- Efficient in Storage given the amount of space taken is less.
- Provides High performance since the indexes are pre-calculated.

Index Storage


◆ How to Bitmap Indexes work?

Fi -- Fact table *i*

Di -- Dimension table *i*

pk -- The primary key column on the dimension table

fk -- The fact table column participating in the join with the dimension tables



```
CREATE BITMAP INDEX bji ON f (d.c1) FROM f, d WHERE d.pk = f.fk
```

```
SELECT SUM(f.sales)
```

```
FROM d, f
```

```
WHERE d.pk = f.fk and d.c1 = 2
```